

# 作业1

姓名：吴双

学号：10164102141

## 作业1

1 软件体系结构

2 详细设计

编写程序

## 1 软件体系结构

Map

要求

班级学生成绩的随机生成

输入：本班同学的学号 输出：<学号, 成绩>

数据

`stuID.csv`文件，每一列为一个学号，在 hdfs 上，shell input:

```
1 | hdfs dfs put stuID.csv input
```

分组MapReduce

要求

求平均成绩：将全班同学每隔5号分为一组，求每组的平均成绩

输入：<学号, 成绩> 输出：<组号, 平均分>

数据

一个 `score.csv`文件，每一列为学号和学生成绩，放入 `HDFS` 中，shell input:

```
1 | hdfs dfs put score.csv input
```

NaturalJoin

要求：

匹配出person中每个人所在的位置信息;

每条记录各个字段之间以空格为分隔符。

数据:

person.txt

address.txt

## Kmeans

数据

k-means.dat

第一行标明K的值和数据个数N, 均为整形, 由","隔开 (如 3,10 表示K=3, N=10)。

之后N行中每行代表一个二维向量, 向量元素均为整形, 由","隔开 (如 1,2 表示向量(1, 2))。

输出:

K行, 每行是一个聚类图心的二维向量, 向量元素均为浮点型 (如 1.1,2.3)。

思路

- 首先需要初始化中心点, 这里使用前四行作为初始中心点, 现将其写入 `cache` 中。
- 每次 `map` 时使用 `setup` 函数读入变量中, 依据所有点与中心店的距离选择其属于的类。
- 在 `reduce` 中, 根据类别进行分组, 对每组聚类重新选择中心点, 将中心点输出到目标文件中。
- 判断两次的中心点是否满足阈值条件, 若不满足, 则将新生成的中心点移动到 `cache` 中, 作为下一次迭代的中心点。
- 迭代结束的标志为: 满足最大迭代次数或满足阈值条件。

## 2 详细设计

### 基础MapReduce

code:

```
1  import org.apache.hadoop.conf.Configuration;
2  import org.apache.hadoop.fs.Path;
3  import org.apache.hadoop.io.IntWritable;
4  import org.apache.hadoop.io.LongWritable;
5  import org.apache.hadoop.io.Text;
6  import org.apache.hadoop.mapreduce.Job;
7  import org.apache.hadoop.mapreduce.Mapper;
8  import org.apache.hadoop.mapreduce.Reducer;
9  import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
10 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
11 import org.apache.hadoop.util.GenericOptionsParser;
12
13 import java.io.IOException;
14 import java.util.Random;
15
16 public class Score {
```

```

17
18     public Score() {
19
20     }
21
22     public static class StuIDMapper
23         extends Mapper<Object, Text, LongWritable, IntWritable> {
24
25         private final static IntWritable one = new IntWritable(1);
26         LongWritable ID = new LongWritable();
27
28         public void map(Object key, Text value, Context context)
29             throws IOException, InterruptedException {
30
31             ID.set(Long.valueOf(value.toString()));
32
33             context.write(this.ID, one);
34
35         }
36     }
37
38     public static class ScoreReducer
39         extends Reducer<LongWritable, IntWritable, LongWritable, IntWritable> {
40         private IntWritable score = new IntWritable();
41         private Random rand = new Random();
42
43         public void reduce(LongWritable ID, Iterable<IntWritable> values, Context
context)
44             throws IOException, InterruptedException {
45             score.set(rand.nextInt(100) + 1);
46             context.write(ID, this.score);
47         }
48     }
49
50
51     public static void main(String[] args) throws Exception {
52
53         org.apache.hadoop.conf.Configuration conf = new Configuration();
54         String[] otherArgs = (new GenericOptionsParser(conf,
args)).getRemainingArgs();
55         if (otherArgs.length < 2) {
56             System.err.println("Usage: wordcount <in>[<in>...] <out>");
57             System.exit(2);
58         }
59
60         Job job = Job.getInstance(conf, "student score");
61         job.setJarByClass(Score.class);
62         job.setMapperClass(Score.StuIDMapper.class);
63         job.setReducerClass(Score.ScoreReducer.class);
64         job.setOutputKeyClass(LongWritable.class);
65         job.setOutputValueClass(IntWritable.class);
66
67         for (int i = 0; i < otherArgs.length - 1; i++) {

```

```

68         FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
69     }
70
71     FileOutputFormat.setOutputPath(job, new Path(otherArgs[otherArgs.length -
72 1]));
73     System.exit(job.waitForCompletion(true) ? 0 : 1);
74 }
75 }

```

## 运行

```

1 | hadoop jar Score.jar input/stuID.csv output
2 | hdfs dfs -ls output/*

```

```

4190558352      38
4190816219      78
4191096730      16
4191634667      33
4192944206     100
4193320729      72
4194369068      29
4195553731      66
4197472805      38
4199736992      91
4200463785      97
4200761214      39
4209095544      30

```

## 分组MapReduce

code:

```

1 | import org.apache.hadoop.conf.Configuration;
2 | import org.apache.hadoop.fs.Path;
3 | import org.apache.hadoop.io.FloatWritable;
4 | import org.apache.hadoop.io.IntWritable;
5 | import org.apache.hadoop.io.LongWritable;
6 | import org.apache.hadoop.io.Text;
7 | import org.apache.hadoop.mapreduce.Job;
8 | import org.apache.hadoop.mapreduce.Mapper;
9 | import org.apache.hadoop.mapreduce.Partitioner;
10 | import org.apache.hadoop.mapreduce.Reducer;
11 | import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
12 | import org.apache.hadoop.mapreduce.lib.input.KeyValueLineRecordReader;
13 | import org.apache.hadoop.mapreduce.lib.input.KeyValueTextInputFormat;
14 | import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 | import org.apache.hadoop.util.GenericOptionsParser;
16 |

```

```

17 import java.io.IOException;
18 import java.util.Random;
19
20 public class GetAvgScore {
21
22     public GetAvgScore() {
23
24     }
25
26     public static class StuScoreMapper
27         extends Mapper<Text, Text, IntWritable, IntWritable> {
28
29         IntWritable Group = new IntWritable();
30         IntWritable Score = new IntWritable();
31
32         public void map(Text key, Text value, Context context)
33             throws IOException, InterruptedException {
34
35             // every 5 students as a group, [1,2,3,4,5] => group1
36             Group.set((Integer.parseInt(key.toString()) - 1) / 5 + 1);
37             Score.set(Integer.parseInt(value.toString()));
38
39             context.write(this.Group, this.Score);
40
41         }
42     }
43
44     public static class AvgScoreReducer
45         extends Reducer<IntWritable, IntWritable, IntWritable, FloatWritable> {
46         private FloatWritable avgscore = new FloatWritable();
47
48         public void reduce(IntWritable Group, Iterable<IntWritable> Score, Context
context)
49             throws IOException, InterruptedException {
50
51             int sum = 0;
52             int count = 0;
53             for (IntWritable val : Score) {
54                 sum += val.get();
55                 count++;
56             }
57             avgscore.set((float) sum / count);
58             context.write(Group, avgscore);
59         }
60
61     }
62
63
64     public static void main(String[] args) throws Exception {
65
66         org.apache.hadoop.conf.Configuration conf = new Configuration();
67         // set seperator
68         conf.set(KeyValueLineRecordReader.KEY_VALUE_SEPERATOR, ",");

```

```

69     String[] otherArgs = (new GenericOptionsParser(conf,
args)).getRemainingArgs();
70     if (otherArgs.length < 2) {
71         System.err.println("Usage: wordcount <in>[<in>...] <out>");
72         System.exit(2);
73     }
74
75     Job job = Job.getInstance(conf, "student avg score");
76     job.setInputFormatClass(KeyValueTextInputFormat.class);
77     job.setJarByClass(GetAvgScore.class);
78
79     job.setMapperClass(GetAvgScore.StuScoreMapper.class);
80     job.setReducerClass(GetAvgScore.AvgScoreReducer.class);
81
82     job.setMapOutputKeyClass(IntWritable.class);
83     job.setMapOutputValueClass(IntWritable.class);
84     job.setOutputKeyClass(IntWritable.class);
85     job.setOutputValueClass(FloatWritable.class);
86
87     for (int i = 0; i < otherArgs.length - 1; i++) {
88         FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
89     }
90
91     FileOutputFormat.setOutputPath(job, new Path(otherArgs[otherArgs.length -
1]));
92     System.exit(job.waitForCompletion(true) ? 0 : 1);
93 }
94
95 }

```

## 运行

```

1 | hadoop jar avgcore.jar input/score.csv output/avgscore/
2 | hdfs dfs -cat output/avgscore/*

```

```

1 | hadoop jar Natural\ join.jar input/person.txt input/address.txt output/natural_join

```

## natural join

code:

```

1 | import org.apache.commons.lang.StringUtils;
2 | import org.apache.hadoop.conf.Configuration;
3 | import org.apache.hadoop.fs.Path;
4 | import org.apache.hadoop.io.Text;
5 | import org.apache.hadoop.mapreduce.Job;

```

```
6 import org.apache.hadoop.mapreduce.Mapper;
7 import org.apache.hadoop.mapreduce.Reducer;
8 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
9 import org.apache.hadoop.mapreduce.lib.input.FileSplit;
10 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
11 import org.apache.hadoop.util.GenericOptionsParser;
12
13
14 import java.io.IOException;
15 import java.util.ArrayList;
16 import java.util.List;
17
18
19 public class natural_join {
20     public natural_join() {
21
22     }
23
24     public static class joinMapper
25         extends Mapper<Object, Text, Text, Text> {
26
27         private static final String PERSON_FLAG = "person";
28         private static final String ADDRESS_FLAG = "address";
29
30         private FileSplit fileSplit;
31         private Text outKey = new Text();
32         private Text outValue = new Text();
33
34
35         public void map(Object key, Text value, Context context)
36             throws IOException, InterruptedException {
37
38             fileSplit = (FileSplit) context.getInputSplit();
39             String filePath = fileSplit.getPath().toString();
40
41             String line = value.toString();
42             String[] fields = StringUtils.split(line, " ");
43
44             // 判断记录来自哪个文件
45             if (filePath.contains(PERSON_FLAG)) {
46                 if (fields.length < 3)
47                     return;
48                 // fields[2] is code
49                 outKey.set(fields[2]);
50                 outValue.set(PERSON_FLAG + "," + line);
51             } else if (filePath.contains(ADDRESS_FLAG)) {
52                 // fields[0] is city code
53                 outKey.set(fields[0]);
54                 outValue.set(ADDRESS_FLAG + "," + fields[1]);
55             }
56
57             context.write(outKey, outValue);
58 }
```

```

59     }
60 }
61 }
62
63 public static class joinReducer
64     extends Reducer<Text, Text, Text, Text> {
65     private static final String PERSON_FLAG = "person";
66     private static final String ADDRESS_FLAG = "address";
67
68     private String fileFlag = null;
69     private String cityName = null;
70
71     private Text outCity = new Text();
72     private Text outPerson = new Text();
73
74     public void reduce(Text key, Iterable<Text> values, Context context)
75         throws IOException, InterruptedException {
76
77         List<String> perosonInfo = new ArrayList<>();
78         for (Text val : values) {
79             String[] fields = StringUtils.split(val.toString(), ",");
80             fileFlag = fields[0];
81             // choose what file it is
82             if (fileFlag.equals(ADDRESS_FLAG)) {
83                 cityName = fields[1];
84                 outCity.set(cityName);
85             } else if (fileFlag.equals(PERSON_FLAG)) {
86                 perosonInfo.add(fields[1]);
87             }
88         }
89
90         // Cartesian product
91         for (String person : perosonInfo) {
92             outPerson.set(person);
93             context.write(outPerson, outCity);
94         }
95     }
96 }
97
98
99 public static void main(String[] args) throws Exception {
100
101     org.apache.hadoop.conf.Configuration conf = new Configuration();
102     conf.set("mapreduce.output.textoutputformat.separator", " ");
103     String[] otherArgs = (new GenericOptionsParser(conf,
104 args)).getRemainingArgs();
105     if (otherArgs.length <= 2) {
106         System.err.println("Usage: natural <in> <in> <out>");
107         System.exit(2);
108     }
109
110     Job job = Job.getInstance(conf, "natural join");
111     job.setJarByClass(natural_join.class);

```



```

111
112     job.setMapperClass(natural_join.joinMapper.class);
113     job.setReducerClass(natural_join.joinReducer.class);
114
115     job.setOutputKeyClass(Text.class);
116     job.setOutputValueClass(Text.class);
117
118     for (int i = 0; i < otherArgs.length - 1; i++) {
119         FileInputFormat.addInputPath(job, new Path(otherArgs[i]));
120     }
121
122     FileOutputFormat.setOutputPath(job, new Path(otherArgs[otherArgs.length -
123 1]));
124     System.exit(job.waitForCompletion(true) ? 0 : 1);
125 }

```

## 运行

```

1 | hadoop jar Natural\ join.jar input/person.txt input/address.txt output/natural_join
2 | hdfs dfs -cat output/natural_join/*

```

```

23 Alva 210000 Nanjing
19 Alger 210000 Nanjing
32 Ansel 210000 Nanjing
1 Aaron 210000 Nanjing
41 Bancroft 210000 Nanjing
27 Andre 210000 Nanjing
62 Bruno 212000 Zhenjiang
59 Boyce 212000 Zhenjiang
46 Barry 213000 Changzhou
37 Atwood 213000 Changzhou
68 Clare 213000 Changzhou
44 Barlow 213000 Changzhou
18 Alfred 214000 Wuxi

```

## Kmeans

code:

class Point:

```

1 | class Point{
2 |     double x;
3 |     double y;
4 |     Point(){
5 |
6 |     }

```

```

7
8     Point(double x,double y){
9         this.x=x;
10        this.y=y;
11    }
12
13    public double EuclideanDis(Point other) {
14        double distance = 0;
15
16        distance = Math.pow((this.x - other.getX()),2) + Math.pow((this.y -
other.getY()),2);
17
18        return Math.sqrt(distance);
19    }
20
21    public double getX()
22    {
23        return x;
24    }
25    public double getY(){
26        return y;
27    }
28 }

```

kmeansMapper:

```

1  import org.apache.commons.lang.StringUtils;
2  import org.apache.hadoop.conf.Configuration;
3  import org.apache.hadoop.fs.FSDataInputStream;
4  import org.apache.hadoop.fs.FileSystem;
5  import org.apache.hadoop.fs.Path;
6  import org.apache.hadoop.io.IntWritable;
7  import org.apache.hadoop.io.LongWritable;
8  import org.apache.hadoop.io.Text;
9  import org.apache.hadoop.mapreduce.Mapper;
10
11  import java.io.BufferedReader;
12  import java.io.FileReader;
13  import java.io.IOException;
14  import java.io.InputStreamReader;
15  import java.net.URI;
16  import java.util.ArrayList;
17  import java.util.List;
18
19  public class kmeansMapper extends Mapper<LongWritable, Text, IntWritable, Text> {
20
21      private List<Point> means;
22
23      /**
24       * reading the data from the distributed cache
25       */
26      public void setup(Context context) throws IOException, InterruptedException {
27          means = new ArrayList<Point>();

```

```

28
29 //
30 //      URI[] cacheFiles = context.getCacheFiles();
31 //      BufferedReader br = new BufferedReader(new
FileReader(cacheFiles[0].toString()));
32
33      Configuration conf = new Configuration();
34      FileSystem fs = FileSystem.get(conf);
35      FSDataInputStream hdfsInStream = fs.open(new Path("output/cache/part-r-
00000"));
36      InputStreamReader isr = new InputStreamReader(hdfsInStream, "utf-8");
37      BufferedReader br = new BufferedReader(isr);
38
39
40      String lineString = null;
41      while((lineString = br.readLine()) != null){
42          String[] keyValue = StringUtils.split(lineString, ",");
43
44          Point tmpCluster = new
Point(Double.parseDouble(keyValue[0]), Double.parseDouble(keyValue[1]));
45          means.add(tmpCluster);
46
47      }
48      br.close();
49  }
50
51  public void map(LongWritable key, Text keyvalue, Context context) throws
IOException, InterruptedException{
52      // ignore first line
53      if (key.get() == 0)
54          return;
55
56      String[] keyValue = StringUtils.split(keyvalue.toString(), ",");
57
58      String x = keyValue[0];
59      String y = keyValue[1];
60
61      Point tmpPoint = new Point(Double.parseDouble(x), Double.parseDouble(y));
62
63      context.write(new IntWritable(findClosest(tmpPoint)), new Text(x + "," + y));
64  }
65
66  /**
67   * method that returns the closest mean from the point
68   * @param value
69   * @return
70   */
71  private int findClosest(Point value){
72      int argmin = 0;
73      double minimalDistance = Double.MAX_VALUE ;
74      for(int i = 0; i<means.size(); i++){
75          Point tmpCluster = means.get(i);
76          double distance = value.EuclideanDis(tmpCluster);

```

```

77         if(distance < minimalDistance){
78             minimalDistance = distance;
79             argmin = i;
80         }
81     }
82     return argmin;
83 }
84 }

```

kmeansReducer:

```

1  import org.apache.commons.lang.StringUtils;
2  import org.apache.hadoop.io.DoubleWritable;
3  import org.apache.hadoop.io.IntWritable;
4  import org.apache.hadoop.io.Text;
5  import org.apache.hadoop.mapreduce.Reducer;
6
7  import java.io.IOException;
8
9  public class kmeansReducer extends Reducer<IntWritable, Text, DoubleWritable,
DoubleWritable> {
10
11     public void reduce(IntWritable key, Iterable<Text> values, Context context)
throws IOException,
12         InterruptedException{
13
14         double sumX = 0.0;
15         double sumY = 0.0;
16         int count = 0;
17         for(Text value : values){
18             String[] keyValue = StringUtils.split(value.toString(),"");
19             sumX += Double.parseDouble(keyValue[0]);
20             sumY += Double.parseDouble(keyValue[1]);
21             count ++;
22         }
23
24
25         context.write(new DoubleWritable(sumX/count), new
DoubleWritable(sumY/count));
26
27     }
28 }

```

kmeansMain:

```

1  import org.apache.commons.lang.StringUtils;
2  import org.apache.hadoop.conf.Configuration;
3  import org.apache.hadoop.fs.FSDataInputStream;
4  import org.apache.hadoop.fs.FSDataOutputStream;
5  import org.apache.hadoop.fs.FileSystem;
6  import org.apache.hadoop.fs.Path;
7  import org.apache.hadoop.io.DoubleWritable;

```

```

8  import org.apache.hadoop.io.IntWritable;
9  import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
11 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
12 import org.apache.hadoop.mapreduce.Job;
13 import org.apache.hadoop.util.GenericOptionsParser;
14
15 import java.io.BufferedReader;
16 import java.io.IOException;
17 import java.io.InputStreamReader;
18 import java.net.URI;
19 import java.net.URISyntaxException;
20
21 public class kmeansMain {
22
23     private static final String CACHED_PATH = "output/cache";
24     private static final String ACTUAL_PATH = "output/means";
25
26     // directory store actual result
27     private static final String CACHED_MEANS = "output/cache/part-r-00000";
28     private static final String ACTUAL_MEANS = "output/means/part-r-00000";
29
30
31     public static void writeFileByline(String dst, String contents) throws
IOException{
32         Configuration conf = new Configuration();
33         Path dstPath = new Path(dst);
34         FileSystem fs = dstPath.getFileSystem(conf);
35         FSDataOutputStream outputStream = null;
36
37         if (!fs.exists(dstPath)) {
38             outputStream = fs.create(dstPath);
39         }else{
40             outputStream = fs.append(dstPath);
41         }
42         contents = contents + "\n";
43         outputStream.write(contents.getBytes("utf-8"));
44         outputStream.close();
45     }
46
47     public static int readFileByLines(String fileName,String meansPath) throws
IOException {
48         Configuration conf = new Configuration();
49         FileSystem fs = FileSystem.get(URI.create(fileName), conf);
50         FSDataInputStream hdfsInStream = fs.open(new Path(fileName));
51         InputStreamReader isr = new InputStreamReader(hdfsInStream, "utf-8");
52         BufferedReader br = new BufferedReader(isr);
53         // get first line k
54         String line = br.readLine();
55         int k = Integer.parseInt(StringUtils.split(line, ",")[0]);
56         int count = 0;
57
58         while ((line = br.readLine()) != null && count < k) {

```

```

59         writeFileByline(meansPath, line);
60         count++;
61     }
62     return k;
63 }
64
65
66 public static void main(String[] args) throws IOException,
67     InterruptedException, ClassNotFoundException, URISyntaxException {
68
69     Configuration conf = new Configuration();
70     String[] otherArgs = (new GenericOptionsParser(conf,
144 args)).getRemainingArgs();
71     if (otherArgs.length < 2) {
72         System.err.println("Usage: kmeans <in> <out>");
73         System.exit(2);
74     }
75
76     int code = 0;
77
78     Path inputPath = new Path(otherArgs[0]);
79     Path outputDir = new Path(otherArgs[1] + "");
80
81     Path cacheMeansPath = new Path(CACHED_MEANS);
82     Path actualMeansPath = new Path(ACTUAL_MEANS);
83
84     Path cachePath = new Path(CACHED_PATH);
85     Path actualPath = new Path(ACTUAL_PATH);
86
87     int k = readFileByLines(otherArgs[0], ACTUAL_MEANS);
88     int maxIterations = 500;
89     double threshold = 0.000001;
90
91
92     // Delete output if exists
93     FileSystem hdfs = FileSystem.get(conf);
94     if (hdfs.exists(outputDir))
95         hdfs.delete(outputDir, true); // recursive delete
96
97     boolean changed = false;
98     int counter = 0;
99
100    while(!changed && counter < maxIterations){
101
102        // Delete output if exists
103        if (hdfs.exists(cachePath))
104            hdfs.delete(cachePath, true);
105        //moving the previous iteration file to the cache directory
106        hdfs.rename(actualPath, cachePath);
107
108        conf.set("threshold", threshold+"");
109        //passing K to the map reduce as a parameter
110        conf.set("k", k+"");

```

```

111         conf.set("mapreduce.output.textoutputformat.separator", ",");
112
113
114         Job kmeans = Job.getInstance(conf, "kmeans "+ (counter + ""));
115
116         // add cache
117         kmeans.addCacheFile(cacheMeansPath.toUri());
118
119         kmeans.setJarByClass(kmeansMapper.class);
120         FileInputFormat.addInputPath(kmeans, inputPath);
121         // set out put path : output/means
122         FileOutputFormat.setOutputPath(kmeans, actualPath);
123
124         kmeans.setMapperClass(kmeansMapper.class);
125         kmeans.setMapOutputKeyClass(IntWritable.class);
126         kmeans.setMapOutputValueClass(Text.class);
127
128         kmeans.setReducerClass(kmeansReducer.class);
129         kmeans.setOutputKeyClass(DoubleWritable.class);
130         kmeans.setOutputValueClass(DoubleWritable.class);
131
132         // Execute job
133         code = kmeans.waitForCompletion(true) ? 0 : 1;
134
135         //checking if the mean is stable
136         BufferedReader file1Reader = new BufferedReader(new
InputStreamReader(hdfs.open(cacheMeansPath)));
137         BufferedReader file2Reader = new BufferedReader(new
InputStreamReader(hdfs.open(actualMeansPath)));
138         for(int i = 0; i<k; i++){
139             String[] keyValue1 = file1Reader.readLine().split(",");
140             String[] keyValue2 = file2Reader.readLine().split(",");
141
142             Point p1 = new
Point(Double.parseDouble(keyValue1[0]),Double.parseDouble(keyValue1[1]));
143             Point p2 = new
Point(Double.parseDouble(keyValue2[0]),Double.parseDouble(keyValue2[1]));
144
145             if(p1.EuclideanDis(p2) <= threshold){
146                 changed = true;
147             }else{
148                 changed = false;
149                 break;
150             }
151         }
152         file1Reader.close();
153         file2Reader.close();
154         counter++;
155         System.out.println("KMEANS finished iteration:>> "+counter + " || means
stable: "+ changed);
156
157     }
158

```

```

159
160         hdfs.rename(actualPath, outputDir);
161
162         System.exit(code);
163
164     }
165
166 }

```

运行

同样的方法，打包成 JAR 包运行，shell input:

```
1 | hadoop jar Kmeans.jar input/k-means.dat output/kmeans
```

output:

```

CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=3106
File Output Format Counters
  Bytes Written=123
KMEANS finished iteration:>> 8 || means stable: true

```

查看运行结果，也就是中心点, shell input:

```
1 | hdfs dfs -cat output/kmeans/*
```

output:

```

98.31,803.89
496.11650485436894,207.3398058252427
493.2277227722772,798.3267326732673
91.05208333333333,206.19791666666666

```